
Exploring Recurrent Architectures for EMG-to-QWERTY Decoding

Akash Madiraju

University of California, Los Angeles
akashm153@ucla.edu

Justin Chen

University of California, Los Angeles
justinc918@ucla.edu

Prabhvir Singh

University of California, Los Angeles
prabab@ucla.edu

Tyler Xiao

University of California, Los Angeles
tylerxiao@ucla.edu

Abstract

This report is prepared for the ECE C147/C247 final project on the emg2qwerty dataset. Our project studies how recurrent sequence models decode typed characters from wrist electromyography (EMG) signals, with a primary focus on comparing vanilla RNN, LSTM, and GRU architectures under the subject 89335547.

1 Introduction

The central goal of this project is to identify and train a model that achieves lower loss and better character error rate (CER) than the base TDS model under the final project guidelines. Because of limited computing resources, all of our experiments were conducted on a single subject, user 89335547. We also kept our models relatively small, targeting fewer than 5 million parameters, both to reduce the risk of overtraining on a limited dataset and to allow faster iteration within the project timeline.

This setting naturally motivated us to explore recurrent architectures. The emg2qwerty task maps continuous wrist sEMG signals to an unsegmented character sequence, so the temporal structure of the signal is central to successful decoding. In contrast to the baseline TDS model, which is primarily convolutional, recurrent neural networks (RNNs) offer a direct mechanism for processing information over time. We therefore focused first on a vanilla RNN as a lightweight temporal baseline, then considered gated recurrent variants such as gated recurrent units (GRUs) and long short-term memory (LSTM) networks to improve optimization and temporal credit assignment on limited data. We also considered a hybrid CNN+RNN architecture as a way to combine spatial feature extraction with temporal modeling if purely recurrent models proved insufficient.

In addition to architecture choice, we considered data augmentation and input reduction strategies to improve generalization without substantially increasing model size. In particular, we explored whether reducing signal detail through lower-frequency resampling or restricting the model to subsets of available channels could encourage the models to rely on more robust patterns rather than subject-specific noise. Together, these design choices frame the main question of this report: which compact recurrent or hybrid architecture most effectively improves decoding quality over the base TDS model on the provided single-user split?

Table 1: Official single-user split used for the main experiments

Split	Sessions	User ID	Notes
Train	16	89335547	Official training sessions from the provided config
Validation	1	89335547	Used for model selection and early comparison
Test	1	89335547	Held out for the final architecture comparison

Table 2: Model families used in the architecture comparison

Model	Encoder configuration	MLP features	Dropout
TDS baseline	4 TDS convolution blocks, kernel width 32	[384]	–
RNN	2-layer bidirectional RNN, hidden size 256	[256]	0.0
GRU	4-layer bidirectional GRU, hidden size 128	[256, 256]	0.1
Hybrid CNN+RNN	4 TDS blocks + 4-layer bidirectional GRU, hidden size 192	[192, 192]	0.1

2 Methods

2.1 Dataset and evaluation setup

This project uses the emg2qwerty dataset released by Meta Reality Labs [2]. For the main comparison, we follow the official single-user split provided by the course codebase for user 89335547. All four model families are trained and evaluated on the same 16/1/1 train/validation/test session split so that differences in performance can be attributed primarily to encoder design rather than to changes in data availability or subject identity.

We evaluate all models primarily with character error rate (CER), which is the main metric emphasized in the project handout and is well matched to the CTC decoding objective. We also track insertion, deletion, and substitution error rates to better diagnose failure modes, since different architectures may over-predict characters, miss characters entirely, or confuse similar temporal patterns in different ways. Validation CER is used for model selection, while the final held-out test session is reserved for the architecture comparison reported in the Results section.

2.2 Model families

Our architecture study compares four encoder families that share the same windowed EMG front end and CTC decoding pipeline. The provided TDSConv model serves as the course baseline and offers a strong convolutional reference point for the task. We then compare three alternatives developed around explicit temporal recurrence: a vanilla bidirectional RNN as the simplest recurrent baseline, a deeper bidirectional GRU to test whether gating improves stability and sequence modeling, and a hybrid CNN+RNN model that applies TDS-style convolutional blocks before a recurrent encoder. This comparison is designed to test whether temporal recurrence alone is sufficient to outperform the baseline or whether a mixed convolutional-recurrent encoder yields a better trade-off on the limited single-user setting.

All four models operate on the same log-spectrogram EMG representation and produce framewise character logits for greedy CTC decoding. This shared output structure lets us compare the effect of encoder choice without changing the loss, decoder, or evaluation metric between runs.

2.3 Training protocol and hyperparameters

Across experiments, we keep the data pipeline fixed and vary only the encoder family and its associated hyperparameters. Input windows have length 8000 with left-context/right-context padding of 1800 and 200 samples, respectively. Raw left and right EMG streams are converted to tensors and transformed into log-spectrogram features using $n_{fft} = 64$ and hop length 16. During training, we apply the augmentations defined in the codebase, including random band rotation, temporal alignment jitter, and SpecAugment. Validation and test runs use the deterministic evaluation pipeline without these stochastic augmentations.

Table 3: Training configurations for the compared model families

Model	LR	Batch size	Epoch budget	Scheduler	Notes
TDS baseline	1e-3	32	40	Warmup + cosine	Original course baseline
RNN	3e-4	32	40	Cosine annealing	2-layer bidirectional recurrent encoder
GRU	1e-4	32	40	Cosine annealing	4-layer gated recurrent encoder
Hybrid CNN+RNN	3e-4	32	40	Cosine annealing	TDS front end + bidirectional GRU encoder

Table 4: Sample architecture comparison with placeholder metrics

Model	Val CER ↓	Test CER ↓	Test IER	Test DER	Test SER
RNN	19.8	27.4	3.1	4.2	20.1
LSTM	18.3	25.6	2.9	3.7	19.0
GRU	17.6	24.9	2.7	3.5	18.8
TDS baseline	21.5	29.3	3.8	4.6	20.9

Optimization follows the PyTorch Lightning and Hydra configuration scaffold used throughout the repository. Unless otherwise noted, runs use batch size 32, four dataloader workers, Adam optimization, and checkpoint selection based on validation CER. Each model uses the CTC greedy decoder for evaluation, which keeps decoding consistent across the convolutional, recurrent, and hybrid architectures. Exact CER and loss values are deferred to the Results section, but the main training settings for the compared models are summarized in Table 3.

In addition to the main architecture comparison, the repository also supports exploratory changes such as reduced channel subsets and related preprocessing variations. Those experiments are treated as secondary analyses rather than the core comparison in this report, so we describe them separately in the Results section only when they contribute directly to the final model discussion.

3 Results

This section should present the main quantitative findings and point the reader to the most important comparisons. A clean structure is to start with the primary architecture comparison on the official single-user split, then follow with one or two ablations such as hyperparameter sensitivity, reduced-data training, or alternative preprocessing.

Table 4 is currently populated with dummy values so the paper has a realistic layout during drafting. In the final version, replace these numbers with actual metrics from the training logs and state which model achieved the best validation CER and which model achieved the best held out test CER. If a model performs inconsistently across validation and test, discuss that explicitly rather than hiding it.

3.1 Potential split and data-efficiency studies

The course handout explicitly suggests studying how much data is needed, how many channels are needed, and how sampling rate affects performance. The table below is included so those additional experiments can be plugged into the final report without restructuring the document.

If these studies are included, the final draft should explain the split protocol precisely. For example, if you sample only a subset of train sessions, state whether the validation and test sessions remain unchanged. If the experiment changes channels or sampling rate, make clear whether the model architecture is otherwise held fixed.

3.2 Qualitative placeholders

If you later generate plots, this report is ready for:

- training and validation CER over epochs,
- CER as a function of hidden size or number of layers,

Table 5: Sample run summary with loss, CER, and compute placeholders

Run ID	Val loss	Val CER	Test loss	Test CER	Time (hr)	Device
A1	0.71	19.8	0.93	27.4	0.9	1×T4
A2	0.65	18.3	0.86	25.6	1.1	1×T4
A3	0.58	17.6	0.81	24.9	1.0	1×T4
A4	0.60	18.1	0.83	25.1	1.4	1×T4
A5	0.56	17.9	0.80	25.0	1.7	1×T4

Table 6: Suggested future split and ablation studies

Study type	Split or subset definition	Primary metric	Purpose
Reduced-data study	25%, 50%, 75%, and 100% of train sessions	CER	Measure data efficiency
Session ablation	First 4, 8, 12, and 16 train sessions	CER	Track how added sessions improve decoding
Channel ablation	Full channel set vs. reduced channel subsets	CER	Estimate minimum useful sensor count
Sampling-rate study	Downsampled EMG features at multiple rates	CER	Measure temporal resolution sensitivity
Cross-user extension	Train on one or more extra users, test per split	CER	Evaluate generalization beyond one user

- CER versus fraction of training data,
- error breakdowns showing insertion, deletion, and substitution trends.

4 Discussion

The discussion should do more than restate the lowest CER. This is the section where you explain what you learned from the comparison. Questions worth answering include: Did gated recurrent models outperform the vanilla RNN? Was the improvement large enough to justify extra compute? Were some runs more stable during training? Did deeper or wider models actually help, or did they mostly overfit the single-user split?

For a strong final write-up, include at least one interpretation tied directly to the course rubric. For example, if the GRU achieves lower CER than the RNN, you might hypothesize that gating helps preserve relevant temporal context in the EMG sequence while avoiding some of the optimization difficulty of a plain recurrent model. If an LSTM does not improve over a GRU, you can discuss whether the added parameters and memory mechanisms were unnecessary for the observed time scale of the task.

4.1 Limitations

This project currently focuses on a single user because the baseline assignment explicitly allows full credit on that setting. However, that choice limits how strongly the results can generalize across users, sessions, and hardware conditions. The final discussion should also note any limits from compute budget, the number of random seeds used, and the extent of hyperparameter search. If error bars or repeated trials are not included, say so directly.

4.2 Reproducibility notes

The final report should include the exact training command, environment, and best checkpoint settings for the main comparison. One compact way to do this is to mention the training command in the discussion or appendix, for example using Hydra overrides such as `user=single_user model=gru_ctc`




Figure 1: Placeholder for a training-curve or ablation plot. Replace this box with a real figure exported from your experiment notebook or analysis script.

Table 7: Compute and reproducibility checklist for final insertion

Item	Placeholder value	Final note to fill in
Framework	PyTorch Lightning	Version and environment file
Main device	1×GPU	Specify model, memory, and provider
Training time	~1 hr / 40 epochs	Replace with measured wall-clock time
Dataset split	Single user 89335547	Mention exact config path
Main command	python -m emg2qwerty.train ...	Insert exact command used

`trainer.max_epochs=40`. Add the exact command, seed, and device information once your final experiment recipe is settled.

Acknowledgments and Disclosure of Funding

This template follows the official NeurIPS 2024 style requested by the course handout and is adapted for the ECE C147/C247 final project on the emg2qwerty dataset. Replace this paragraph with any acknowledgments, compute-credit notes, or dataset access notes that you want in the final submission.

References

- [1] Alex Graves, Santiago Fernandez, Faustino Gomez, and Jurgen Schmidhuber. Connectionist temporal classification: Labelling unsegmented sequence data with recurrent neural networks. In *Proceedings of the 23rd International Conference on Machine Learning*, pages 369–376, 2006.
- [2] Viswanath Sivakumar, Jeffrey Seely, Alan Du, Sean R. Bittner, Adam Berenzweig, Anuoluwapo Bolarinwa, Alexandre Gramfort, and Michael I. Mandel. *emg2qwerty: A large dataset with baselines for touch typing using surface electromyography*, 2024. arXiv:2410.20081.
- [3] ECE C147/C247 staff. *ECE C147/C247, Winter 2026 project guidelines*, 2026.

A Appendix / supplemental material

Use the appendix for any material that supports the main claims but does not fit cleanly into the seven-page class report body. Good candidates include extra ablation plots, confusion analyses, exact hyperparameter grids, full reproduction commands, and additional tables for multi-user or cross-session experiments.

NeurIPS Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer [Yes], [No], or [NA].
- [NA] means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for NA).

The checklist answers are an integral part of your paper submission. They are visible to the reviewers, area chairs, senior area chairs, and ethics reviewers. You will be asked to also include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While "[Yes]" is generally preferable to "[No]", it is perfectly acceptable to answer "[No]" provided a proper justification is given (e.g., "error bars are not reported because it would be too computationally expensive" or "we were unable to find the license for the dataset we used"). In general, answering "[No]" or "[NA]" is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- **Delete this instruction block, but keep the section heading “NeurIPS paper checklist”,**
- **Keep the checklist subsection headings, questions/answers and guidelines below.**
- **Do not modify the questions and only use the provided macros for your answers.**

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory Assumptions and Proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental Result Reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer NA means that the paper does not include experiments.

- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).

- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental Setting/Details

Question: Does the paper specify all the training and test details (e.g., data splits, hyper-parameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment Statistical Significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments Compute Resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.

- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code Of Ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader Impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer NA means that the paper poses no such risks.

- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New Assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and Research with Human Subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional Review Board (IRB) Approvals or Equivalent for Research with Human Subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.